

Graves 2013, “Generating Sequences with Recurrent Neural Networks”

Johannes Bausch and Jack Kamm

14 November 2017

- 1 RNN and LSTM
- 2 Generating text
- 3 Generating handwriting

Table of Contents

- 1 RNN and LSTM
- 2 Generating text
- 3 Generating handwriting

RNN




Figure 1: **Deep recurrent neural network prediction architecture.** The circles represent network layers, the solid lines represent weighted connections and the dashed lines represent predictions.

$$h_t^1 = \mathcal{H}(W_{ih^1}x_t + W_{h^1h^1}h_{t-1}^1 + b_h^1)$$

$$h_t^n = \mathcal{H}(W_{ih^n}x_t + W_{h^{n-1}h^n}h_t^{n-1} + W_{h^nh^n}h_{t-1}^n + b_h^n)$$

$$\hat{y}_t = b_y + \sum_{i=1}^N W_{h^ny}h_t^n$$

$$y_t = \mathcal{Y}(\hat{y}_t)$$

Likelihood of input, loss function:

$$\mathbb{P}(\mathbf{x}) = \prod_{t=1}^T \mathbb{P}(x_{t+1} \mid y_t)$$

$$\mathcal{L}(\mathbf{x}) = - \sum_{t=1}^T \log \mathbb{P}(x_{t+1} \mid y_t)$$

LSTM




Figure 2: Long Short-term Memory Cell


Instead of

$$h_t^n = \mathcal{H}(W_{ih^n}x_t + W_{h^{n-1}h^n}h_t^{n-1} + W_{h^n h^n}h_{t-1}^n + b_h^n)$$


do we have instead the following?

$$h_t^n, c_t^n = \mathcal{H}(x_t, c_t^{n-1}, h_t^{n-1}, h_{t-1}^n)$$


For a really nice explanation of LSTM see:
colah.github.io/posts/2015-08-Understanding-LSTMs




Gate=sigmoid +
multiplication



0='nothing thru'
1='all thru'




Information flows along cell state



$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$


Forget gate controls what to keep from previous



$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$


$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

Input gate controls what to add from input



$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

Update cell state from forget and input gates



$$o_t = \sigma(W_o [h_{t-1}, x_t] + b_o)$$

$$h_t = o_t * \tanh(C_t)$$

Output gate controls what to output from cell state

Backpropagation thru time

- Train with “backpropagation thru time”
 - Turn head sideways and do backpropagation
 - i.e. reverse-mode automatic differentiation
- Truncated backpropagation thru time
 - Improve numerical stability by truncating gradients if they get too big as we go backwards thru the sequence

Reference other papers?

Table of Contents

- 1 RNN and LSTM
- 2 Generating text
- 3 Generating handwriting

Text Prediction

Basic framework for text prediction

$$y_t = \mathbb{P}(x_{t+1})$$

Per-character vs per-word prediction

Penn Treebank Test Set

Penn Treebank
Perplexity? BPC?
Regularization schemes

Wikipedia Experiments

- Wikipedia experiments
- `karpathy.github.io/2015/05/21/rnn-effectiveness` has some nice visualizations for understanding what's going on, e.g. what the inner neurons represent

Table of Contents

- 1 RNN and LSTM
- 2 Generating text
- 3 Generating handwriting

Handwriting experiments

- Handwriting experiments
- Mixture density outputs
 - Basic framework for real valued outputs
- Figure 10 is cool

Handwriting Synthesis

- Handwriting synthesis
- Input, output sequences have different lengths
- Biased vs unbiased vs primed sampling

Synthesis Network




Figure: $\phi_t(h_t^1) = \text{distr over positions}$; $w_t = \mathbf{c}\phi_t = \text{distr over characters}$